Commodity Derivatives Pricing Engine

Theoretical Foundations and Implementation

Schwartz-Smith Two-Factor Model · Kalman Filter Calibration
SABR Stochastic Volatility · Black-76 Pricing
Historical VaR/CVaR · Stress Testing

Author	**Pablo M. Paniagua**
Project	QCEX — Quantum Commodity Exchange
Commodities	Cacao (ICE CC1) · Gas Natural TTF · Uranio U3O8
References	Schwartz (1997) · Schwartz & Smith (2000) Hagan et al. (2002) · Black (1976)
Stack	Python 3.11 · NumPy · SciPy · Streamlit · Plotly

Contents

1 Introduction and Motivation **3**
1.1 Choice of Commodities . . . 3
1.2 Architectural Overview . . . 3

2 Theoretical Background: Commodity Price Dynamics **3**
2.1 The Theory of Storage . . . 3
2.1.1 Convenience Yield . . . 4
2.2 Stylised Facts of Commodity Prices . . . 4

3 The Schwartz-Smith Two-Factor Model **4**
3.1 Motivation for Two Factors . . . 4
3.2 Physical Measure Dynamics . . . 5
3.2.1 Factor Properties . . . 5
3.3 Risk-Neutral Measure . . . 5
3.4 Futures Pricing Formula . . . 6
3.4.1 Key Structural Properties of $A(\tau)$ . . . 6
3.5 Convenience Yield Extraction . . . 6
3.6 Exact Discretisation and State-Space Form . . . 7

4 Kalman Filter and Maximum Likelihood Calibration **7**
4.1 Kalman Filter Recursion . . . 7
4.1.1 Initialisation . . . 7
4.1.2 Prediction Step . . . 7
4.1.3 Innovation . . . 8
4.1.4 Update Step (Kalman Gain) . . . 8
4.2 Log-Likelihood and MLE Calibration . . . 8
4.2.1 Identification . . . 8
4.3 Model Diagnostics . . . 8

5 Monte Carlo Simulation **8**
5.1 Path Simulation Algorithm . . . 8
5.2 Monte Carlo Option Pricing . . . 9

6 Black-76 Model for Options on Futures **9**
6.1 Assumptions . . . 9
6.2 Pricing Formulas . . . 9
6.3 The Greeks . . . 10
6.3.1 Implied Volatility Inversion . . . 10

7 SABR Stochastic Volatility Model **10**
7.1 Motivation . . . 10
7.2 Dynamics . . . 10
7.3 Hagan (2002) Implied Volatility Approximation . . . 11
7.3.1 ATM Limit ($F = K$) . . . 11
7.4 Smile Interpretation . . . 12
7.4.1 Skew . . . 12
7.4.2 Smile Curvature . . . 12
7.5 SABR Calibration . . . 12
7.5.1 Calibration Results (Cacao, $T = 6M$) . . . 12
7.6 Vol Surface and Term Structure . . . 12

7.6.1 SABR Limitations . . . 13

8 Risk Analytics **13**
8.1 Value at Risk . . . 13
8.2 Stress Testing . . . 13
8.3 Volatility Cone . . . 14

9 System Architecture and Implementation **14**
9.1 Module Structure . . . 14
9.2 Data Architecture . . . 14
9.3 Calibration Performance . . . 14
9.4 Key Design Decisions . . . 14
9.4.1 Why Schwartz-Smith vs Single-Factor? . . . 14
9.4.2 Why $\beta = 0.5$ in SABR? . . . 15
9.4.3 Why Historical VaR over Parametric? . . . 15

10 Commodity-Specific Analysis **15**
10.1 Cacao (Cocoa) . . . 15
10.1.1 Market Structure . . . 15
10.1.2 ENSO and Seasonal Dynamics . . . 15
10.1.3 Calibrated Parameters . . . 15
10.2 Gas Natural TTF . . . 15
10.2.1 Post-2022 Regime . . . 15
10.2.2 Seasonality . . . 16
10.2.3 Fast Mean-Reversion . . . 16
10.3 Uranio U3O8 . . . 16
10.3.1 OTC Market . . . 16
10.3.2 Slow Mean-Reversion . . . 16

11 Weather Indicators and Integration into the Pricing Model **16**
11.1 Motivation: Weather as a Supply and Demand Driver . . . 16
11.2 Cocoa: ENSO and West African Rainfall . . . 17
11.2.1 Transmission Mechanism . . . 17
11.2.2 Indicators Implemented . . . 17
11.2.3 Integration into the Pricing Model . . . 17
11.3 Natural Gas TTF: Heating Degree Days (HDD) . . . 18
11.3.1 Transmission Mechanism . . . 18
11.3.2 Heating Degree Days (HDD) . . . 18
11.3.3 Sinusoidal Fallback . . . 18
11.4 Uranium U3O8: No Direct Weather Driver . . . 18
11.5 Module Architecture: `data/weather.py` . . . 19

12 References **19**

1. Introduction and Motivation

This document provides the complete theoretical derivation underlying the **Commodity Derivatives Pricing Engine**, a production-grade quantitative system for pricing, calibrating, and risk-managing derivatives on physical commodities. The project is intentionally structured around a set of commodities that are analytically rich but underrepresented in standard quant portfolios.

1.1. Choice of Commodities

Commodity	Exchange / Proxy	Unit	Key analytical feature
Cacao (Cocoa)	ICE Futures CC1	USD/MT	Supply concentration, ENSO correlation, crisis
Gas Natural TTF	ICE TTF / Henry Hub NG	EUR/MWh	Stationarity seasonal, geopolitical basis risk, LN
Uranio U3O8	OTC / Sprott URA ETF	USD/lb	OTC market, nuclear policy driven, thin smile

Table 1: Commodities covered and their analytical characteristics

1.2. Architectural Overview

The system is structured in four modular layers:

1. **Data layer** (`data/fetcher.py`): yfinance as primary source with synthetic mean-reverting fallback, ensuring the app always runs.
2. **Models layer** (`models/`): Schwartz-Smith two-factor model with Kalman Filter; SABR stochastic vol with Hagan approximation; Black-76 option pricer with full Greeks.
3. **Analytics layer** (`analytics/risk.py`): Historical and parametric VaR/CVaR, vol cone, stress testing, drawdown analysis.
4. **Dashboard layer** (`app/`): Four-page Streamlit application with Plotly visualisations, live data feeds, and interactive calibration.

2. Theoretical Background: Commodity Price Dynamics

2.1. The Theory of Storage

The fundamental relationship between spot and futures prices in commodity markets was formalised by **Working (1949)** through the *theory of storage*. The cost-of-carry model states:

$$F(t,T) = S_t \cdot e^{(r+u-y)(T-t)} \tag{1}$$

where:

- S_t is the spot price at time t
- r is the risk-free interest rate
- u is the cost of physical storage (annualised fraction)
- y is the **convenience yield** — the implicit benefit of holding the physical commodity
- $T - t = \tau$ is the time to maturity

2.1.1. Convenience Yield

The convenience yield y is not directly observable. It is extracted by *bootstrapping* from the observed futures curve:

$$y(t,T) = r + u - \frac{1}{\tau}\ln\left(\frac{F(t,T)}{S_t}\right) \tag{2}$$

Definition 2.1 (Market Structure). *When $y > r + u$, the futures curve is in **backwardation** ($F < S$), indicating tight physical supply: agents pay a premium for immediate delivery. When $y < r + u$, the market is in **contango** ($F > S$), implying storage costs dominate.*

Remark 2.1. *For commodities with strong seasonality (e.g. gas, cacao), the convenience yield $y(\tau)$ itself has term structure and is not constant. This motivates the stochastic models in Section 3.*

2.2. Stylised Facts of Commodity Prices

Commodity prices exhibit several empirical features that distinguish them from equity prices:

F1. **Mean reversion**: Prices tend toward a long-run equilibrium determined by production costs. Deviations induced by supply shocks are transient.

F2. **Seasonality**: Demand and supply are periodic (weather-driven for gas, harvest-driven for cacao).

F3. **Inventory effects**: Low inventories amplify price volatility and strengthen backwardation (Fama and French, 1987).

F4. **Negative skewness**: Supply disruptions cause sharp upward spikes. The return distribution has heavier left tails than normal when conditioned on negative outcomes for producers.

F5. **Stochastic volatility**: Implied vol smiles in commodity options are pronounced, requiring beyond-Black-Scholes models.

F6. **Imperfect futures correlation**: Near-term and long-term futures prices are not perfectly correlated, invalidating single-factor models.

3. The Schwartz-Smith Two-Factor Model

3.1. Motivation for Two Factors

Schwartz (1997) proposed a one-factor mean-reverting model for commodity prices. However, Schwartz and Smith (2000) demonstrated empirically that a single latent factor implies perfect correlation across the entire futures term structure, a hypothesis strongly rejected by the data. The two-factor model resolves this by separating the price process into:

- A **short-term factor** χ_t: transient deviations from equilibrium (weather shocks, inventory surprises, geopolitical events)
- A **long-term factor** ξ_t: the equilibrium price level, evolving slowly with production costs, energy transition, and long-run supply/demand

$$\ln S_t = \chi_t + \xi_t \tag{3}$$

3.2. Physical Measure Dynamics

Under the real-world probability measure $\mathbb{P}$:

$$\mathrm{d}\chi_t = -\kappa\,\chi_t\,\mathrm{d}t + \sigma_\chi\,\mathrm{d}W_\chi^{\mathbb{P}} \tag{4}$$

$$\mathrm{d}\xi_t = \mu_\xi\,\mathrm{d}t + \sigma_\xi\,\mathrm{d}W_\xi^{\mathbb{P}} \tag{5}$$

$$\mathrm{d}W_\chi^{\mathbb{P}}\,\mathrm{d}W_\xi^{\mathbb{P}} = \rho\,\mathrm{d}t \tag{6}$$

Parameter	Description	Interpretation
$\kappa > 0$	Mean-reversion speed	Rate at which χ_t returns to zero. Half-life $= \ln 2/\kappa$
μ_ξ	Long-term drift	Expected appreciation of equilibrium price
σ_χ	Short-term vol	Volatility of transient shocks
σ_ξ	Long-term vol	Volatility of equilibrium level
$\rho \in (-1, 1)$	Correlation	Between short-term and long-term Brownian motions

Table 2: Schwartz-Smith model parameters under $\mathbb{P}$

3.2.1. Factor Properties

Equation (4) is an **Ornstein-Uhlenbeck** (OU) process. By Itô's formula:

$$\chi_t = \chi_0 e^{-\kappa t} + \sigma_\chi \int_0^t e^{-\kappa(t-s)}\,\mathrm{d}W_\chi^{\mathbb{P}}(s) \tag{7}$$

which implies:

$$\mathbb{E}[\chi_t] = \chi_0\,e^{-\kappa t} \xrightarrow{t\to\infty} 0 \tag{8}$$

$$\mathrm{Var}[\chi_t] = \frac{\sigma_\chi^2}{2\kappa}\left(1 - e^{-2\kappa t}\right) \xrightarrow{t\to\infty} \frac{\sigma_\chi^2}{2\kappa} \tag{9}$$

The **half-life** of the short-term factor is $h = \ln 2/\kappa$. For our calibrated commodities:

- **Cacao** ($\kappa = 0.80$): half-life $\approx$ 10.4 months
- **Gas TTF** ($\kappa = 2.10$): half-life $\approx$ 4.0 months
- **Uranium** ($\kappa = 0.30$): half-life $\approx$ 27.6 months

Equation (5) is a **Brownian motion with drift**, the long-term permanent component. Unlike χ_t, it does not mean-revert.

3.3. Risk-Neutral Measure

Introducing constant market prices of risk λ_χ (short-term) and λ_ξ (long-term), by Girsanov's theorem:

$$\mathrm{d}W_\chi^{\mathbb{Q}} = \mathrm{d}W_\chi^{\mathbb{P}} + \frac{\lambda_\chi}{\sigma_\chi}\,\mathrm{d}t \tag{10}$$

$$\mathrm{d}W_\xi^{\mathbb{Q}} = \mathrm{d}W_\xi^{\mathbb{P}} + \frac{\lambda_\xi}{\sigma_\xi}\,\mathrm{d}t \tag{11}$$

Under $\mathbb{Q}$ the dynamics become:

$$\mathrm{d}\chi_t = \left(-\kappa\chi_t - \lambda_\chi\right)\,\mathrm{d}t + \sigma_\chi\,\mathrm{d}W_\chi^{\mathbb{Q}} \tag{12}$$

$$\mathrm{d}\xi_t = \left(\mu_\xi - \lambda_\xi\right)\,\mathrm{d}t + \sigma_\xi\,\mathrm{d}W_\xi^{\mathbb{Q}} \tag{13}$$

The short-term factor now mean-reverts to $-\lambda_\chi/\kappa$ under $\mathbb{Q}$, encoding the market price of short-term supply risk.

3.4. Futures Pricing Formula

Theorem 3.1 (Schwartz-Smith Futures Price)**.** *Under the two-factor model, the time-t futures price for delivery at T is:*

$$\ln F(t,T) = e^{-\kappa\tau}\chi_t + \xi_t + A(\tau) \tag{14}$$

where $\tau = T - t$ and the deterministic correction is:

$$A(\tau) = \underbrace{(\mu_\xi - \lambda_\xi)\tau}_{\textit{risk-adj. drift}} - \underbrace{\frac{1-e^{-\kappa\tau}}{\kappa}\lambda_\chi}_{\textit{S-T risk prem.}} + \underbrace{\frac{\sigma_\xi^2\tau}{2}}_{\textit{Jensen (L-T)}} + \underbrace{\frac{\sigma_\chi^2(1-e^{-2\kappa\tau})}{4\kappa}}_{\textit{Jensen (S-T)}} + \underbrace{\frac{\rho\sigma_\chi\sigma_\xi(1-e^{-\kappa\tau})}{\kappa}}_{\textit{cross-factor}} \tag{15}$$

Sketch. By the no-arbitrage condition, $F(t,T) = \mathbb{E}^{\mathbb{Q}}[S_T \mid \mathcal{F}_t]$. Since $\ln S_T = \chi_T + \xi_T$ is Gaussian under $\mathbb{Q}$ (both factors follow linear SDEs), we apply the lognormal moment generating function:

$$\mathbb{E}^{\mathbb{Q}}[e^X] = e^{\mathbb{E}^{\mathbb{Q}}[X] + \frac{1}{2}\operatorname{Var}^{\mathbb{Q}}[X]}$$

Computing $\mathbb{E}^{\mathbb{Q}}[\chi_T \mid \mathcal{F}_t]$, $\mathbb{E}^{\mathbb{Q}}[\xi_T \mid \mathcal{F}_t]$, and the conditional variances and covariance yields (14)–(15). □

3.4.1. Key Structural Properties of $A(\tau)$

As $\tau \to \infty$: $e^{-\kappa\tau} \to 0$, so $\ln F(t,T) \to \xi_t + A(\tau)$ — long-dated futures depend *only* on the permanent factor ξ_t.

As $\tau \to 0$: $F(t,T) \to S_t$ — no-arbitrage boundary condition.

The slope of the log-futures curve at short maturities is:

$$\left.\frac{\partial \ln F}{\partial \tau}\right|_{\tau=0} = -\kappa\chi_t + (\mu_\xi - \lambda_\xi) - \lambda_\chi + \sigma_\xi^2/2 + \rho\sigma_\chi\sigma_\xi$$

A large negative χ_t (supply shock pushing prices above equilibrium) induces steep backwardation in the short end of the curve.

3.5. Convenience Yield Extraction

From the model, the implied convenience yield consistent with (14) is:

$$y(\tau) = r + u - \frac{A(\tau)}{\tau} - \frac{e^{-\kappa\tau}-1}{\tau}\chi_t + \frac{\xi_t - \ln S_t}{\tau} \tag{16}$$

In the implementation, we use the simpler observable formula (2) bootstrapped from market futures prices.

3.6. Exact Discretisation and State-Space Form

Let Δt be the observation interval. Define the state vector $\mathbf{x}_t = (\chi_t, \xi_t)^\top$. The exact discretisation of the joint SDE system is:

$$\mathbf{x}_t = \mathbf{F}\,\mathbf{x}_{t-1} + \boldsymbol{\eta}_t, \qquad \boldsymbol{\eta}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}) \tag{17}$$

$$\mathbf{F} = \begin{pmatrix} e^{-\kappa\Delta t} & 0 \\ 0 & 1 \end{pmatrix}, \qquad \mathbf{Q} = \begin{pmatrix} Q_{11} & Q_{12} \\ Q_{12} & Q_{22} \end{pmatrix} \tag{18}$$

$$Q_{11} = \frac{\sigma_\chi^2}{2\kappa}\left(1 - e^{-2\kappa\Delta t}\right) \tag{19}$$

$$Q_{22} = \sigma_\xi^2\,\Delta t \tag{20}$$

$$Q_{12} = \frac{\rho\sigma_\chi\sigma_\xi}{\kappa}\left(1 - e^{-\kappa\Delta t}\right) \tag{21}$$

The observation equation for n futures contracts with maturities $\tau_1, \ldots, \tau_n$ is:

$$\mathbf{y}_t = \mathbf{H}\,\mathbf{x}_t + \mathbf{d} + \boldsymbol{\varepsilon}_t, \qquad \boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{R}) \tag{22}$$

where $\mathbf{y}_t = (\ln F(t, T_1), \ldots, \ln F(t, T_n))^\top$ and:

$$\mathbf{H} = \begin{pmatrix} e^{-\kappa\tau_1} & 1 \\ \vdots & \vdots \\ e^{-\kappa\tau_n} & 1 \end{pmatrix}, \qquad \mathbf{d} = \begin{pmatrix} A(\tau_1) \\ \vdots \\ A(\tau_n) \end{pmatrix}, \qquad \mathbf{R} = s^2\mathbf{I}_n \tag{23}$$

with s^2 a homoscedastic measurement noise variance.

4. Kalman Filter and Maximum Likelihood Calibration

4.1. Kalman Filter Recursion

The linear Gaussian state-space system (17)–(22) admits the exact **Kalman Filter** for optimal (minimum mean-squared error) estimation of latent states $\{\chi_t, \xi_t\}$ from observable log-futures prices.

4.1.1. Initialisation

The short-term factor χ_t has a known stationary distribution:

$$\chi_0 \sim \mathcal{N}\left(0, \frac{\sigma_\chi^2}{2\kappa}\right)$$

The long-term factor ξ_t has a diffuse (uninformative) prior with large variance.

4.1.2. Prediction Step

$$\hat{\mathbf{x}}_{t|t-1} = \mathbf{F}\,\hat{\mathbf{x}}_{t-1|t-1} \tag{24}$$

$$\mathbf{P}_{t|t-1} = \mathbf{F}\,\mathbf{P}_{t-1|t-1}\,\mathbf{F}^\top + \mathbf{Q} \tag{25}$$

4.1.3. Innovation

$$\mathbf{v}_t = \mathbf{y}_t - \mathbf{H}\,\hat{\mathbf{x}}_{t|t-1} - \mathbf{d} \tag{26}$$

$$\mathbf{S}_t = \mathbf{H}\,\mathbf{P}_{t|t-1}\,\mathbf{H}^\top + \mathbf{R} \tag{27}$$

4.1.4. Update Step (Kalman Gain)

$$\mathbf{K}_t = \mathbf{P}_{t|t-1}\,\mathbf{H}^\top\,\mathbf{S}_t^{-1} \tag{28}$$

$$\hat{\mathbf{x}}_{t|t} = \hat{\mathbf{x}}_{t|t-1} + \mathbf{K}_t\,\mathbf{v}_t \tag{29}$$

$$\mathbf{P}_{t|t} = (\mathbf{I} - \mathbf{K}_t\,\mathbf{H})\,\mathbf{P}_{t|t-1} \tag{30}$$

4.2. Log-Likelihood and MLE Calibration

The innovations $\mathbf{v}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{S}_t)$ (under a well-specified model). The **prediction error decomposition** of the log-likelihood is:

$$\ell(\boldsymbol{\theta}) = -\frac{1}{2}\sum_{t=1}^{T}\Big[\ln|\mathbf{S}_t| + \mathbf{v}_t^\top\,\mathbf{S}_t^{-1}\,\mathbf{v}_t + n\ln(2\pi)\Big] \tag{31}$$

where $\boldsymbol{\theta} = (\kappa, \mu_\xi, \sigma_\chi, \sigma_\xi, \rho, \lambda_\chi, \lambda_\xi, s)$ is the parameter vector.
Maximum likelihood estimates are obtained by solving:

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg\max_{\boldsymbol{\theta}}\ \ell(\boldsymbol{\theta})$$

via **L-BFGS-B** (gradient-based, bounded) with multiple random restarts to mitigate local optima.

4.2.1. Identification

Parameters μ_ξ and λ_ξ are not separately identified from cross-sectional futures data alone (only their difference $\mu_\xi - \lambda_\xi$ enters $A(\tau)$). Identification requires either time-series variation in the level of the forward curve, or imposing restrictions on the market price of risk.

4.3. Model Diagnostics

A correctly specified model produces *whitened innovations*:

- $\mathbb{E}[\mathbf{v}_t] \approx \mathbf{0}$: no systematic pricing bias
- $\text{AC}(\mathbf{v}_t, \mathbf{v}_{t-k}) \approx 0$ for $k \geq 1$: no serial correlation (Ljung-Box test)
- Approximate normality (Jarque-Bera test)

5. Monte Carlo Simulation

5.1. Path Simulation Algorithm

Given calibrated parameters $\hat{\boldsymbol{\theta}}$ and filtered states $(\hat{\chi}_T, \hat{\xi}_T)$ at the last observation date, future spot paths are simulated under $\mathbb{P}$:

Step 1. Compute Cholesky decomposition of the 2×2 correlation matrix:

$$\mathbf{\Sigma} = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} = \mathbf{L}\mathbf{L}^\top$$

Step 2. At each step $t = 1, \ldots, N$, draw correlated increments:

$$(Z_\chi, Z_\xi)^\top = \mathbf{L}\,\mathbf{Z}, \quad \mathbf{Z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_2)$$

Step 3. Update factors using exact discretisation for χ and Euler-Maruyama for ξ:

$$\chi_t = \chi_{t-1}\, e^{-\kappa \Delta t} + \sigma_\chi \sqrt{\frac{1 - e^{-2\kappa\Delta t}}{2\kappa}}\, Z_\chi \tag{32}$$

$$\xi_t = \xi_{t-1} + \mu_\xi \Delta t + \sigma_\xi \sqrt{\Delta t}\, Z_\xi \tag{33}$$

Step 4. Reconstruct spot: $S_t = \exp(\chi_t + \xi_t)$

The exact discretisation for χ_t (step 3) avoids discretisation bias that would occur with Euler-Maruyama on the OU process.

5.2. Monte Carlo Option Pricing

For a European call on spot with strike K, maturity T, and risk-free rate r:

$$C_{\text{MC}} = e^{-rT} \frac{1}{N_{\text{paths}}} \sum_{i=1}^{N_{\text{paths}}} \max(S_T^{(i)} - K,\ 0) \tag{34}$$

The standard error is:

$$\widehat{\text{SE}} = e^{-rT} \frac{\hat{\sigma}_{\text{payoff}}}{\sqrt{N_{\text{paths}}}}$$

yielding a 95% confidence interval of width $2 \times 1.96\,\widehat{\text{SE}}$.

6. Black-76 Model for Options on Futures

6.1. Assumptions

Black (1976) extends the Black-Scholes framework to options whose underlying is a futures contract $F(t, T_F)$ rather than a spot asset. The key simplification: futures require no initial investment, so no cost-of-carry term appears.

Under $\mathbb{Q}$, the futures price follows:

$$\mathrm{d}F = \sigma F\ \mathrm{d}W^{\mathbb{Q}}$$

6.2. Pricing Formulas

$$C = e^{-r\tau}\left[F\,N(d_1) - K\,N(d_2)\right] \tag{35}$$

$$P = e^{-r\tau}\left[K\,N(-d_2) - F\,N(-d_1)\right] \tag{36}$$

$$d_1 = \frac{\ln(F/K) + \frac{1}{2}\sigma^2\tau}{\sigma\sqrt{\tau}}, \qquad d_2 = d_1 - \sigma\sqrt{\tau} \tag{37}$$

where $N(\cdot)$ is the standard normal CDF.

6.3. The Greeks

Greek	Formula (call)	Financial meaning	Units
Δ	$e^{-r\tau}N(d_1)$	Sensitivity to F	Dimensionless
Γ	$\dfrac{e^{-r\tau}\phi(d_1)}{F\sigma\sqrt{\tau}}$	Convexity w.r.t. F	per unit of F
ν (Vega)	$F\,e^{-r\tau}\phi(d_1)\sqrt{\tau}$	Sensitivity to σ	per unit of vol
Θ	$-\dfrac{F\,e^{-r\tau}\phi(d_1)\sigma}{2\sqrt{\tau}} - rKe^{-r\tau}N(d_2)$	Time decay (annual)	per year
Second-order Greeks (computed via finite differences in the implementation)			
Vanna	$\partial^2 V/\partial F\,\partial\sigma$	Mixed F-vol sensitivity	Crucial for barrier/digital opt
Volga	$\partial^2 V/\partial\sigma^2$	Vol convexity; captures smile risk	Vol smile curvature premium

Table 3: Black-76 Greeks for call options. $\phi(\cdot)$ denotes the standard normal PDF.

6.3.1. Implied Volatility Inversion

Given a market option price C_{mkt}, the implied vol $\hat{\sigma}$ satisfies:

$$C_{\text{B76}}(F, K, r, \hat{\sigma}, \tau) = C_{\text{mkt}}$$

Inverted via **Newton-Raphson**: starting from $\sigma_0 = 0.30$, iterate:

$$\sigma_{n+1} = \sigma_n - \frac{C_{\text{B76}}(\sigma_n) - C_{\text{mkt}}}{\nu(\sigma_n)}$$

convergence guaranteed in $O(10)$ iterations due to the monotonicity of C in σ and the convexity of the vega function.

7. SABR Stochastic Volatility Model

7.1. Motivation

Black-76 assumes constant volatility, producing a *flat vol surface*: all strikes and maturities are priced with the same σ. In reality, commodity options exhibit:

- **Negative skew**: OTM puts are more expensive than OTM calls (supply risk)
- **Smile curvature**: Both wings trade above ATM vol (high vol-of-vol)
- **Term structure of smile**: The shape evolves with maturity

The **SABR model** (Hagan, Kumar, Lesniewski and Woodward, 2002) addresses this with a two-dimensional stochastic process for the forward F and its volatility σ.

7.2. Dynamics

$$\mathrm{d}F = \sigma F^{\beta}\ \mathrm{d}W_1 \tag{38}$$
$$\mathrm{d}\sigma = \alpha\,\sigma\ \mathrm{d}W_2 \tag{39}$$
$$\mathrm{d}W_1\ \mathrm{d}W_2 = \rho\ \mathrm{d}t \tag{40}$$

Param.	Role	Commodity intuition
$\alpha > 0$	Initial vol level	Scales the smile vertically; calibrates to ATM vol
$\beta \in [0,1]$	CEV elasticity	$\beta = 0$: normal; $\beta = 1$: lognormal; $\beta = 0.5$: standard for commodities
$\rho \in (-1,1)$	Skew driver	$\rho < 0$: negative skew (puts expensive). Supply risk in cacao, gas
$\nu > 0$	Vol of vol	Smile curvature. High ν in energy (spikes); lower in OTC uranium

Table 4: SABR model parameters and their economic interpretation

Remark 7.1. *The SABR model has no mean reversion in σ, unlike the Heston model. This is intentional: it prioritises tractability of the smile approximation over long-term vol dynamics.*

7.3. Hagan (2002) Implied Volatility Approximation

The exact SABR model has no closed-form option price. Hagan et al. derived a **singular perturbation expansion** in the vol-of-vol parameter α, yielding:

$$\sigma_B(K,T) \approx \frac{\alpha}{D(K)} \cdot \frac{z}{\chi(z)} \cdot C(T) \tag{41}$$

where:

$$D(K) = (FK)^{(1-\beta)/2}\left[1 + \frac{(1-\beta)^2}{24}\ln^2\frac{F}{K} + \frac{(1-\beta)^4}{1920}\ln^4\frac{F}{K}\right] \tag{42}$$
$$z = \frac{\nu}{\alpha}(FK)^{(1-\beta)/2}\ln\frac{F}{K} \tag{43}$$
$$\chi(z) = \ln\left[\frac{\sqrt{1-2\rho z+z^2}+z-\rho}{1-\rho}\right] \tag{44}$$
$$C(T) = 1 + T\left[\frac{(1-\beta)^2\alpha^2}{24(FK)^{1-\beta}} + \frac{\rho\beta\nu\alpha}{4(FK)^{(1-\beta)/2}} + \frac{\nu^2(2-3\rho^2)}{24}\right] \tag{45}$$

7.3.1. ATM Limit ($F = K$)

Taking the limit $K \to F$ (L'Hôpital's rule gives $z/\chi(z) \to 1$):

$$\sigma_{\mathrm{ATM}}(T) \approx \frac{\alpha}{F^{1-\beta}}\left[1 + T\left(\frac{(1-\beta)^2\alpha^2}{24F^{2(1-\beta)}} + \frac{\rho\beta\nu\alpha}{4F^{1-\beta}} + \frac{\nu^2(2-3\rho^2)}{24}\right)\right] \tag{46}$$

At $T = 0$: $\sigma_{\mathrm{ATM}} \approx \alpha/F^{1-\beta}$, so α is directly the ATM vol rescaled by the forward.

7.4. Smile Interpretation

7.4.1. Skew

The first-order slope of the smile around ATM is controlled by ρ:

$$\left.\frac{\partial \sigma_B}{\partial K}\right|_{K=F} \approx -\frac{\rho\nu}{2F^{(1-\beta)/2}}$$

Negative ρ induces negative skew (downward-sloping smile from left to right): the market assigns higher implied vol to downside strikes (OTM puts). This is the dominant effect for commodities with supply concentration risk (cacao, gas in winter).

7.4.2. Smile Curvature

The second-order curvature ("smile" or "butterfly") is driven by ν:

$$\left.\frac{\partial^2 \sigma_B}{\partial K^2}\right|_{K=F} \propto \nu^2$$

High ν (as in gas TTF: $\nu = 0.85$) produces pronounced wings. Low ν (uranium: $\nu = 0.42$) produces a flatter, more symmetric profile.

7.5. SABR Calibration

With β fixed (typically $\beta = 0.5$ for commodities), the remaining three parameters (α, ρ, ν) are calibrated per maturity slice by **weighted least squares** on implied vol:

$$(\hat{\alpha}, \hat{\rho}, \hat{\nu}) = \arg\min_{\alpha,\rho,\nu} \sum_{i=1}^{n} w_i \left[\sigma_B(K_i, T; \alpha, \rho, \nu) - \sigma_i^{\text{mkt}}\right]^2 \tag{47}$$

Solved via multi-start **L-BFGS-B**. Calibration is performed independently for each maturity slice; the resulting parameter surface $\{(\hat{\alpha}(T), \hat{\rho}(T), \hat{\nu}(T))\}$ describes the full vol surface.

7.5.1. Calibration Results (Cacao, $T = 6M$)

	α	ρ	ν	**RMSE (bp)**
Prior (calibrated)	0.52	−0.45	0.68	—
Fitted (synthetic data)	0.80	−0.50	0.56	33.6

Table 5: SABR calibration to synthetic market smile (11 strikes, T=6M, $\beta = 0.5$)

7.6. Vol Surface and Term Structure

The full **implied vol surface** is the map $(K, T) \mapsto \sigma_B(K, T)$ computed from (41) for a grid of strikes $K \in [0.70F, 1.30F]$ and maturities $T \in \{1M, 2M, \ldots, 2Y\}$.

Key cross-sectional features:

- **Term structure of ATM vol**: $\sigma_{\text{ATM}}(T)$ typically increases with T ("normal" vol term structure) but can invert during supply crises (backwardation of vol)
- **Skew term structure**: The skew $|\rho|$ tends to decrease with maturity as long-term supply becomes more elastic
- **Smile decays**: At very long maturities, $C(T)$ corrections become large and the Hagan approximation may fail (arbitrage in wings for $T > 3Y$)

7.6.1. SABR Limitations

L1. Absorbed probability mass: The SABR process can hit zero in finite time (when $\beta < 1$). For long maturities, the approximation (41) can violate calendar spread arbitrage.

L2. Static per-slice: Calibrating independently per slice does not ensure a consistent dynamic model for volatility. Extensions include SABR-LSV (Local Stochastic Volatility).

L3. Thin markets: For uranium (OTC), the smile data is sparse and noisy. The calibration is more uncertain; ν estimation requires wide bid-ask spreads.

8. Risk Analytics

8.1. Value at Risk

Definition 8.1 (Historical VaR). *The $(1-p)$-quantile of the empirical distribution of scaled log-returns:*

$$VaR_{p,h} = -Q_p\left(\{r_i\sqrt{h}\}_{i=1}^{T}\right) \tag{48}$$

where $r_i = \ln(S_i/S_{i-1})$ are daily log-returns and h is the holding period in days.

Definition 8.2 (Parametric (Gaussian) VaR).

$$VaR_{p,h}^{Gauss} = \hat{\sigma}\,\Phi^{-1}(1-p)\,\sqrt{h} \tag{49}$$

where $\hat{\sigma}$ is the sample standard deviation of daily log-returns.

Definition 8.3 (Expected Shortfall (CVaR)).

$$CVaR_{p,h} = -\mathbb{E}\left[r\sqrt{h} \;\middle|\; r\sqrt{h} < -VaR_{p,h}\right] \tag{50}$$

CVaR is a coherent risk measure (Artzner et al., 1999) and the regulatory standard under FRTB (Basel III/IV).

Remark 8.1. *Commodity returns exhibit excess kurtosis (fat tails), so Gaussian VaR typically* underestimates *true risk. Historical VaR is preferred for commodities.*

8.2. Stress Testing

The stress scenarios implemented are commodity-specific, calibrated to historical crisis events:

Commodity	Scenario	Shock (%)	Mechanism
Cacao	El Niño severo	-45	Drought in West Africa; crop failure
Cacao	Fungal disease	-35	Swollen shoot destroys 30% harvest
Gas TTF	Invierno frío extremo	$+80$	HDD +40%; heating demand surge
Gas TTF	Resolución conflicto	-50	Russian pipeline reopening; supply glut
Uranio	Kazatomprom recorte	$+50$	Kazakhstan cuts output 20% (43% global share)
Uranio	Accidente nuclear	-40	Level 4–5 event; reactor shutdowns

Table 6: Selected stress scenarios (full list in `analytics/risk.py`)

8.3. Volatility Cone

The **volatility cone** provides a distributional view of realised volatility across measurement horizons:

$$\hat{\sigma}_w(t) = \sqrt{252}\,\hat{\sigma}\big(\{r_i\}_{i=t-w+1}^{t}\big) \tag{51}$$

For each window $w \in \{5, 10, 21, 63, 126, 252\}$, we compute the empirical percentiles $\{P_{10}, P_{25}, P_{50}, P_{75}, P_{90}\}$ and the current value. This identifies whether current vol is historically cheap or expensive — a key input to options strategies.

9. System Architecture and Implementation

9.1. Module Structure

Module	Key functions	Dependencies
`models/schwartz_smith.py`	`kalman_filter`, `calibrate`, `simulate_paths`	NumPy, SciPy
`models/sabr.py`	`SABRModel`, `sabr_implied_vol`, `sabr_greeks`	NumPy, SciPy
`analytics/risk.py`	`historical_var`, `stress_table`, `vol_cone`	NumPy, Pandas
`data/fetcher.py`	`fetch_prices`, `build_term_structure`	yfinance, Pandas
`app/main.py`	Streamlit entry point	Streamlit, Plotly

Table 7: Module structure and key components

9.2. Data Architecture

The data layer uses a **two-tier fallback** strategy:

T1. Primary: `yfinance` REST API (CC=F cacao, NG=F gas, URA uranium ETF)

T2. Fallback: Synthetic data generated via the calibrated Schwartz 1-factor OU process, ensuring the app always runs in offline or demo environments

All data fetches are cached at the Streamlit layer with a 1-hour TTL.

9.3. Calibration Performance

9.4. Key Design Decisions

9.4.1. Why Schwartz-Smith vs Single-Factor?

Single-factor models produce perfect cross-sectional correlation in the futures curve: a property strongly rejected empirically. Schwartz-Smith (2000) tested this on crude oil, copper, and gold, finding significant two-factor structure. For cacao (where weather-driven short-term spikes are largely uncorrelated with the long-run equilibrium set by production costs), the two-factor decomposition is particularly valuable.

Model	**Time (s)**	**RMSE**	**Method**
Kalman Filter (single pass, $T = 260$, $n = 6$)	< 0.01	—	Exact recursion
Schwartz-Smith MLE (5 restarts)	≈ 48	—	L-BFGS-B
SABR calibration (11 strikes)	0.20	33.6 bp	Multi-start L-BFGS-B
Black-76 implied vol (Newton-Raphson)	< 0.001	$< 10^{-8}$	10 iterations
Monte Carlo option price (50,000 paths)	≈ 2	± 0.0004	Cholesky corr.

Table 8: Computational performance benchmarks

9.4.2. Why $\beta = 0.5$ in SABR?

The CEV exponent β controls the relationship between the forward level and its instantaneous vol. Pure log-normal ($\beta = 1$) is standard for equity and FX, but it implies volatility grows linearly with price — problematic for commodities that can spike to extreme levels. Normal ($\beta = 0$) implies negative prices are possible. $\beta = 0.5$ (geometric Brownian motion with stochastic vol) is the industry standard for commodity options desks: it keeps prices positive while allowing moderate vol-level dependence.

9.4.3. Why Historical VaR over Parametric?

Commodity returns have excess kurtosis and negative skewness, violating the normality assumption of parametric VaR. For cacao, the 2023–24 crisis produced daily returns in the $\pm 8\%$ range, far outside a Gaussian $\pm 3\sigma$ band. Historical VaR captures this empirically without distributional assumptions.

10. Commodity-Specific Analysis

10.1. Cacao (Cocoa)

10.1.1. Market Structure

The global cocoa market is characterised by extreme geographic concentration: Ivory Coast and Ghana together account for approximately 65% of world production, creating a structurally supply-inelastic commodity vulnerable to weather shocks.

10.1.2. ENSO and Seasonal Dynamics

El Niño Southern Oscillation (ENSO) episodes correlate strongly with reduced rainfall in West Africa. The 2023–24 El Niño triggered a crop failure that drove spot prices from $\approx \$3,000/\text{MT}$ to over $\$12,000/\text{MT}$ — a historical anomaly that produced the high calibrated volatility $\sigma_\chi = 0.38$ and the strong negative SABR skew $\rho = -0.45$.

10.1.3. Calibrated Parameters

10.2. Gas Natural TTF

10.2.1. Post-2022 Regime

Following the energy crisis triggered by the reduction of Russian pipeline supplies to Europe in 2022, TTF gas volatility shifted to a structurally higher regime. Historical vol moved from

Parameter	Value (calibrated)
κ (mean-reversion)	0.80 $\text{yr}^{-1} \Rightarrow$ half-life $\approx$ 10.4 months
σ_χ (S-T vol)	38% p.a. (crisis-calibrated)
σ_ξ (L-T vol)	18% p.a.
ρ (correlation)	-0.25
SABR ρ (skew)	-0.45 (strong supply risk skew)
SABR ν (smile)	0.68

$\approx$ 30% to peaks above 200% annualised, with the SABR vol-of-vol parameter $\nu = 0.85$ reflecting persistent spike risk.

10.2.2. Seasonality

Gas demand is temperature-driven (Heating Degree Days, HDD). The seasonal index is modelled as:

$$\text{Seasonal}(m) = 1 + 0.35 \cos\left(\frac{(m-1) \cdot 2\pi}{12}\right)$$

with $m \in \{1, \ldots, 12\}$ indexing months. This generates the winter premium embedded in the forward curve.

10.2.3. Fast Mean-Reversion

The high $\kappa = 2.10$ yr^{-1} (half-life $\approx$ 4 months) reflects the fact that gas price spikes are rapidly corrected as LNG imports substitute for pipeline shortfalls. The term structure of gas futures is typically steep in backwardation during winter and flattens by summer, consistent with the model.

10.3. Uranio U3O8

10.3.1. OTC Market

Uranium is not exchange-traded. Prices are set via bilateral long-term contracts (utilities and miners) and a thin spot market, leading to low liquidity and wide bid-ask spreads. The Sprott Physical Uranium Trust (URA ETF) is used as a market proxy.

10.3.2. Slow Mean-Reversion

With $\kappa = 0.30$ yr^{-1} (half-life $\approx$ 28 months), uranium prices respond slowly to supply/demand shocks. The 2021–24 price appreciation from $\approx$ \$30/lb to $>$ \$100/lb reflects a multi-year structural shift driven by: (a) renewed interest in nuclear power post-2022 energy crisis, (b) Kazatomprom production cuts (Kazakhstan supplies 43% of global primary production), and (c) financial demand via physical uranium trusts.

11. Weather Indicators and Integration into the Pricing Model

11.1. Motivation: Weather as a Supply and Demand Driver

Physical commodity prices are directly conditioned by meteorological variables. Unlike financial assets, whose supply and demand are primarily determined by macroeconomic variables, physical commodities have production and consumption cycles with strong climatic dependence.

QCEX integrates real weather data from the **Open-Meteo** free API (`archive-api.open-meteo.com`), which provides daily historical temperature and precipitation data without requiring an API key. When the connection is unavailable, the system generates synthetic data using sinusoidal functions calibrated to the historical seasonal patterns of each geographic zone.

11.2. Cocoa: ENSO and West African Rainfall

11.2.1. Transmission Mechanism

World cocoa production is concentrated in a narrow tropical belt. Ivory Coast and Ghana together produce 65% of global supply, making the market structurally vulnerable to any meteorological disruption in that region.

The main climate driver is the **ENSO** (*El Niño-Southern Oscillation*) phenomenon. El Niño episodes displace the intertropical convergence zone, reducing precipitation in West Africa during the critical months of flowering and pod formation (February–April). The causal chain is:

El Niño → Reduced rainfall in Ghana/Ivory Coast → Water stress in plantations → Lower yield per tree → Reduced supply → Price increase

The 2023–24 episode perfectly illustrates this mechanism: the most severe El Niño since 1998 caused a 25% production decline and drove prices from \$3,000/MT to over \$12,000/MT in less than twelve months.

11.2.2. Indicators Implemented

The `data/weather.py` module computes the following indicators for the reference location (Abidjan, 5.35°N / 4.00°W):

Indicator	**Calculation**	**Bullish signal for price**
Precipitation 30d	Mean daily rainfall (mm)	$<$ historical mean
Dry days (30d)	Days with rainfall < 0.5 mm	> 15 days
Temperature anomaly	$\bar{T}_{30d} - \bar{T}_{\text{hist}}$	$> +1.5$°C
ENSO proxy	$(\bar{P}_{30d} - \bar{P}_{\text{hist}})/\bar{P}_{\text{hist}}$	< -0.20

Table 9: Weather indicators for Cocoa

11.2.3. Integration into the Pricing Model

Meteorological indicators can be incorporated as regressors in the long-term factor ξ_t of the Schwartz-Smith model. The proposed extension is:

$$d\xi_t = \left(\mu_\xi + \beta_{\text{ENSO}} \cdot \text{ENSO}_t\right) dt + \sigma_\xi \, dW_\xi \tag{52}$$

where ENSO_t is the proxy index calculated from the precipitation anomaly. A $\beta_{\text{ENSO}} < 0$ implies that dry periods (negative ENSO) generate upward pressure on the equilibrium price level.

11.3. Natural Gas TTF: Heating Degree Days (HDD)

11.3.1. Transmission Mechanism

Natural gas is the only commodity in the system with a dominant meteorological **demand** driver (unlike cocoa, whose driver is supply-side). Heating demand in Europe is directly correlated with outdoor temperature. The reference location is Berlin (52.52°N / 13.40°E), representative of the Central European market that determines TTF prices.

11.3.2. Heating Degree Days (HDD)

Heating Degree Days are the energy industry standard metric for quantifying thermal demand:

$$\mathrm{HDD}_d = \max\big(T_{\text{base}} - T_d,\ 0\big) \tag{53}$$

where $T_{\text{base}} = 15.5°\text{C}$ is the temperature below which heating activates, and T_d is the daily mean temperature. The 30-day cumulative HDD:

$$\mathrm{HDD}_{30} = \sum_{d=t-29}^{t} \max(15.5 - T_d,\ 0) \tag{54}$$

The anomaly relative to the historical mean determines whether winter is colder or warmer than usual, generating respectively upward or downward pressure on TTF.

11.3.3. Sinusoidal Fallback

When real data is unavailable, the system approximates daily temperature in Berlin with:

$$T(m) = 12 + 10\cos\left(\frac{(m-1)\cdot 2\pi}{12}\right) + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0,4) \tag{55}$$

This sinusoidal function captures the annual cycle: minimum temperature in January ($T \approx 2°\text{C}$) and maximum in July ($T \approx 22°\text{C}$), with daily volatility of 2°C. **This is not a forecast** — it is a fallback to ensure the application functions without an internet connection.

11.4. Uranium U3O8: No Direct Weather Driver

Uranium is the only commodity in the system without direct weather dependence. Its prices are driven by:

1. **Nuclear policy**: number of reactors in operation and under construction
2. **Supply concentration**: Kazakhstan (Kazatomprom) controls 43% of global production
3. **Financial demand**: physical funds such as the Sprott Uranium Trust
4. **Long-term contract cycles**: utilities sign 5–10 year supply contracts, decoupling spot prices from short-term equilibrium

The uranium indicators panel in QCEX replaces meteorological variables with energy policy and market structure indicators, which are the real price drivers.

11.5. Module Architecture: `data/weather.py`

The module follows a two-tier architecture:

T1. Real data: HTTP request to the Open-Meteo API with a 5-second timeout. If the response is valid, indicators are computed on real historical data. The panel displays the `OPEN-METEO` label.

T2. Synthetic fallback: if the request fails (no internet, timeout, etc.), data is generated via the seasonal stochastic process described above. The panel displays the `SYNTHETIC` label.

This design ensures the application **always runs**, both in connected environments (real data) and in offline demonstrations (synthetic data).

12. References

[1] Artzner, P., Delbaen, F., Eber, J.-M. & Heath, D. (1999). Coherent measures of risk. *Mathematical Finance*, 9(3), 203–228.

[2] Black, F. (1976). The pricing of commodity contracts. *Journal of Financial Economics*, 3(1–2), 167–179.

[3] Fama, E. & French, K. (1987). Commodity futures prices: Some evidence on forecast power, premiums, and the theory of storage. *Journal of Business*, 60(1), 55–73.

[4] Hagan, P., Kumar, D., Lesniewski, A. & Woodward, D. (2002). Managing smile risk. *Wilmott Magazine*, September, 84–108.

[5] Kalman, R.E. (1960). A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1), 35–45.

[6] Schwartz, E.S. (1997). The stochastic behavior of commodity prices: Implications for valuation and hedging. *Journal of Finance*, 52(3), 923–973.

[7] Schwartz, E.S. & Smith, J.E. (2000). Short-term variations and long-term dynamics in commodity prices. *Management Science*, 46(7), 893–911.

[8] Working, H. (1949). The theory of price of storage. *American Economic Review*, 39(6), 1254–1262.